

```
In [1]: # Imports
import warnings
warnings.filterwarnings("ignore")
from sklearn.datasets import load_boston
from random import seed
from random import randrange
from csv import reader
from math import sqrt
from sklearn import preprocessing
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from prettytable import PrettyTable
from sklearn.linear_model import SGDRegressor
from sklearn import preprocessing
from sklearn.metrics import mean_squared_error
```

## sklearn Boston Dataset

```
In [2]: boston = load_boston()
```

```
In [3]: boston.keys()
```

```
Out[3]: dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename'])
```

```
In [4]: boston.data.shape
```

```
Out[4]: (506, 13)
```

506 data samples, with 13 features each

```
In [5]: boston.feature_names
```

```
Out[5]: array(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD',
               'TAX', 'PTRATIO', 'B', 'LSTAT'], dtype='<U7')
```

```
In [6]: print (boston.DESCR)
```

```
.. _boston_dataset:
```

Boston house prices dataset

-----

**\*\*Data Set Characteristics:\*\***

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.

:Attribute Information (in order):

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B  $1000(B_k - 0.63)^2$  where  $B_k$  is the proportion of blacks by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

.. topic:: References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In

Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

```
In [7]: df = pd.DataFrame(boston.data)
```

```
In [8]: df.head(5)
```

Out[8]:

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [9]: # sklearn Train Test split
from sklearn.model_selection import train_test_split
```

```
In [10]: X_train, X_test, Y_train, Y_test = train_test_split(boston.data, boston.target,
, test_size=0.3, random_state=4)
```

```
In [11]: # Standardizing the Train and Test Data
# Perform fit on Train data and Transform on both Train and Test Data
scaler = preprocessing.StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [12]: # To visualize how the data looks post standardization, lets convert it into a
pandas dataframe
standardized_df = pd.DataFrame(X_train)
standardized_df.head(5)
```

Out[12]:

	0	1	2	3	4	5	6	7	
0	-0.425469	-0.470768	-0.954686	-0.231455	-0.919581	0.215100	-0.747410	0.454022	-0.764441
1	-0.426323	2.992576	-1.330157	-0.231455	-1.227311	-0.883652	-1.691588	3.163428	-0.651511
2	-0.385190	-0.470768	-0.705828	4.320494	-0.423795	-0.125423	0.818985	-0.353904	-0.199911
3	-0.249268	-0.470768	-0.423497	-0.231455	-0.158805	-0.228336	1.021567	-0.021755	-0.651511
4	-0.365945	0.395068	-1.030363	-0.231455	0.157472	3.102729	-0.060078	-0.646202	-0.538611

```
In [13]: standardized_df['house_price'] = Y_train  
standardized_df.head(5)
```

Out[13]:

	0	1	2	3	4	5	6	7	
0	-0.425469	-0.470768	-0.954686	-0.231455	-0.919581	0.215100	-0.747410	0.454022	-0.76441
1	-0.426323	2.992576	-1.330157	-0.231455	-1.227311	-0.883652	-1.691588	3.163428	-0.65151
2	-0.385190	-0.470768	-0.705828	4.320494	-0.423795	-0.125423	0.818985	-0.353904	-0.19991
3	-0.249268	-0.470768	-0.423497	-0.231455	-0.158805	-0.228336	1.021567	-0.021755	-0.65151
4	-0.365945	0.395068	-1.030363	-0.231455	0.157472	3.102729	-0.060078	-0.646202	-0.53861

```
In [14]: # Retrieving Random 5 samples from dataset  
standardized_df.sample(5)
```

Out[14]:

	0	1	2	3	4	5	6	7	
352	-0.421647	-0.470768	-1.018720	-0.231455	-0.398151	-0.587019	0.001420	-0.515405	-0.538
205	-0.168268	-0.470768	1.241383	-0.231455	2.687694	-1.605559	0.909423	-1.062438	-0.538
191	-0.155219	-0.470768	1.241383	-0.231455	2.687694	-1.475404	0.916658	-0.965939	-0.538
35	-0.407931	0.481652	-0.755308	-0.231455	-1.073446	-0.989592	0.283589	1.935812	-0.312
59	0.443181	-0.470768	1.025997	-0.231455	1.337103	0.081918	0.544051	-0.477515	1.606

```
In [15]: standardized_df.sample(5)
```

Out[15]:

	0	1	2	3	4	5	6	7	
111	0.111574	-0.470768	1.025997	-0.231455	1.337103	0.687291	0.771956	-0.564122	1.606
75	-0.392820	-0.470768	1.577561	-0.231455	0.576327	0.319527	1.093918	-0.787639	-0.651
250	-0.424806	-0.470768	0.259046	-0.231455	-1.022158	-0.022509	-0.541211	0.563541	-0.538
13	0.225702	-0.470768	1.025997	-0.231455	-0.210094	-0.007375	-0.143281	-0.178357	1.606
106	2.055917	-0.470768	1.025997	-0.231455	1.567900	-0.649070	0.858777	-0.905370	1.606

The image shows a video player interface with a dark background. At the top left, there is a logo with the letter 'A' and the text 'Implement SGD for linear regression'. To the right of the logo are two buttons: 'Watch later' and 'Share'. The main content of the video is handwritten mathematical equations in white and orange ink. The first equation is the cost function: 
$$\min_{w, b} \sum_{i=1}^n (y_i - w^T x_i - b)^2 = \mathcal{L}(w, b)$$
 Below this, there are two partial derivative equations, each preceded by a checkmark. The first derivative is with respect to the weights  $w$ : 
$$\checkmark \quad \frac{\partial \mathcal{L}(w)}{\partial w} = \left( \sum_{i=1}^n (-2x_i) (y_i - w^T x_i - b) \right)$$
 The second derivative is with respect to the bias  $b$ : 
$$\checkmark \quad \frac{\partial \mathcal{L}}{\partial b} = \left( \sum (-2) (y_i - w^T x_i - b) \right)$$
 The expressions in parentheses for both derivatives are circled in orange. A play button is visible in the center of the video frame.

From the optimization function above, we need to compute weights ( $W$ ) and Bias ( $B$ ) for the linear regression model. For a two feature input system, the model is a line equation. In the current problem, we have 13 features, so the model is a hyper plane

We need to compute partial derivatives with respect to Weights and Bias for the optimization problem

Implement SGD for linear regression

iterative algo

$$\vec{w}_{j+1} = \vec{w}_j - \gamma \left( \frac{\partial L}{\partial \vec{w}} \right) \vec{w}_j$$

$$b_{j+1} = b_j - \gamma \left( \frac{\partial L}{\partial b} \right) b_j$$

(100 iterations)

k-iterations

$\vec{w}_k \approx \vec{w}_{k+1}$

$b_k \approx b_{k+1}$

reduce  $\gamma$  with each iter<sup>n</sup>.

$\gamma_0 = 1$

iter 0  $\rightarrow \gamma = 1$

1  $\rightarrow \gamma = \gamma/2$

2  $\rightarrow \gamma = \gamma/2$

...

MORE VIDEOS

APPLIED COURSE

4:19 / 5:53

YouTube

```

In [68]: def _customSGD(dataset, learning_rate, n_epochs, batch_size, initialize_random
        _coefs=False, variable_learning=False):
        """
        # Stochastic Gradient Descent
        -- Gradient Descent computes the cost gradient for all the samples involve
        d in the dataset
            where as the SGD computes cost gradient of 1 sample at each iteration.
        -- Instead of one sample at a time, if a batch_size is provided it then be
        comes Batch GD

        Expectations:
            dataset - a pandas dataframe
            learning_rate - an intezer or float
            n_epochs - an intezer
            batch_Size - an intezer

        """

        # Raise TypeErrors when required inputs are of not desired format

        no_of_training_sample, no_of_features = dataset.shape
        no_of_features -= 1 # subtracing the labels column from features

        if initialize_random_coefs:
            # Random initialization
            bias = np.random.rand(1)
            weights = np.random.rand(no_of_features)
        else:
            # Initialization with zeros
            bias = 0
            weights = np.zeros(shape=(1,no_of_features),dtype="double")

        # for future add tqdm function hear to know the progress of epochs
        for _epoch in range(n_epochs):
            # updating Learning rate for every 100 epochs
            if variable_learning == True:
                if _epoch%100 == 0:
                    learning_rate /= 2.0

            # Local variables for needing to predict and compute error
            _b, _w, _paderivateW, _paderivateB = bias, weights, np.zeros(shape=(
1,no_of_features)), 0

            # since this going to be mini batch SGD
            miniBatch = dataset.sample(batch_size)

            # reference - https://stackoverflow.com/questions/40144769/how-to-sel
            ect-the-last-column-of-dataframe
            ylabels = np.array(miniBatch.iloc[:,-1]) # labels column
            xfeatures = np.array(miniBatch.drop(miniBatch.columns[green(miniBatch.co
lums)-1], axis=1))

            for index in range(batch_size):
                # Refer first image for formulae
                # partial derivative w.r.t Weights
                #  $dL/dw = \text{Summation}(-2x * (y - (wTx+b)))$ 

```



```

        _paderivateW += (-2) * (xfeatures[index]) * (ylabels[index] - (np
.dot(_w,xfeatures[index]) + _b))

        # partial derivate w.r.t bias
        # dL/db = Summation(-2 * (y - (wTx+b)))
        _paderivateB += (-2) * (ylabels[index] - (np.dot(_w, xfeatures[in
dex]) + _b))

        # Updating the weights, bias at the end of epoch
        # refer second image for formulae
        bias = (_b - learning_rate*(_paderivateB)/batch_size)
        weights = (_w - learning_rate*(_paderivateW)/batch_size)

        # print (bias, weights, f"At the end of epoch {_epoch} out of {n_epoch
s}")

    return weights, bias

```

```
In [69]: fweights, fbias = _customSGD(standardized_df, 0.01, 750, 25) # fixed learning
rate
```

```
In [70]: print (f"Weights\n===== \n{fweights}")
```

```

Weights
=====
[[-0.89519037  1.12956652 -0.35030211  1.13923212 -1.35240868  2.366946
 -0.23322673 -3.29670507  2.01243158 -1.33301936 -1.6670795  0.73886308
 -3.70065531]]

```

```
In [71]: print (f"Bias\n===== \n{fbias}")
```

```

Bias
=====
[22.19157247]

```

```
In [82]: # variable learning rate
vweights, vbias = _customSGD(standardized_df, 0.01, 750, 25, False, True) # fi
xed learning rate
```

```
In [83]: print (f"Weights\n===== \n{vweights}")
```

```

Weights
=====
[[-0.70785641  0.42531237 -0.56623238  1.08724266 -0.43664337  2.76976657
 -0.21268316 -1.31358926  0.44666002 -0.37973611 -1.54300038  0.76709466
 -2.62892823]]

```

```
In [84]: print (f"Bias\n===== \n{vbias}")
```


```

Bias
=====
[19.11943868]

```


```
In [72]: # since we are done with the training lets do inference
# with fixed learning rate
predicted_house_yhat_list = []
for index in range(len(X_test)):
    yhat = np.dot(fweights, X_test[index]) + fbias
    predicted_house_yhat_list.append(np.asscalar(yhat)) # converting the yhat
    prediction into a scalar value
    # >>> np.asscalar(np.array([24]))
    # 24
```

```
In [73]: # Lets plot Ground Truth Price v/s Predictions
plt.figure(figsize=(10,10))
plt.title("Ground Truth Price v/s Prediction Price with Fixed Learning Rate")
plt.scatter(Y_test, predicted_house_yhat_list)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.grid()
plt.show()
```



```
In [85]: # with Variable Learning rate
vpredicted_house_yhat_list = []
for index in range(len(X_test)):
    yhat = np.dot(vweights, X_test[index]) + vbias
    vpredicted_house_yhat_list.append(np.asscalar(yhat)) # converting the yhat
prediction into a scalar value
# >>> np.asscalar(np.array([24]))
# 24
```

```
In [86]: # Lets plot Ground Truth Price v/s Predictions
plt.figure(figsize=(10,10))
plt.title("Ground Truth Price v/s Prediction Price with variable Learning Rate")
plt.scatter(Y_test, vpredicted_house_yhat_list)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.grid()
plt.show()
```



```
In [87]: # Mean Squared Error - Fixed LR, Variable LR
MSE_flr = mean_squared_error(Y_test, predicted_house_yhat_list)
MSE_vlr = mean_squared_error(Y_test, vpredicted_house_yhat_list)

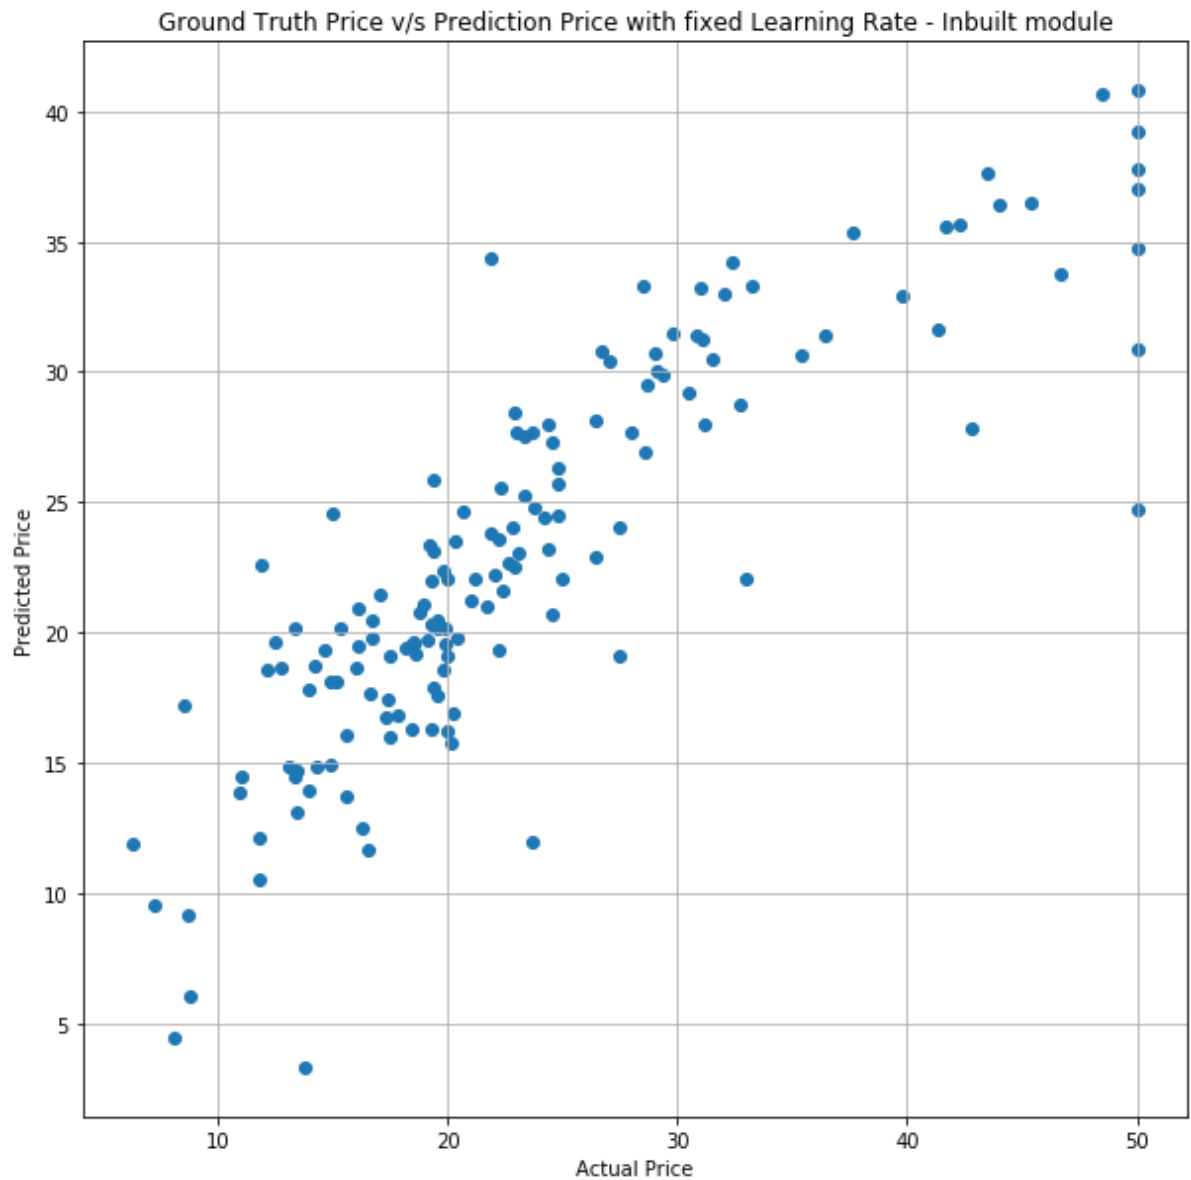
print (f"Fixed LR MSE {MSE_flr}")
print (f"Variable LR MSE {MSE_vlr}")
```

Fixed LR MSE 30.64473411784173  
Variable LR MSE 42.52533567999639

```
In [75]: # Lets perform the same with inbuilt module of sklearn

sgd = SGDRegressor(learning_rate="constant", eta0=0.01, penalty=None, max_iter=50)
sgd.fit(X_train,Y_train)
y_hat_sgd=sgd.predict(X_test)
```

```
In [76]: # Lets plot Ground Truth Price v/s Predictions
plt.figure(figsize=(10,10))
plt.title("Ground Truth Price v/s Prediction Price with fixed Learning Rate -
Inbuilt module")
plt.scatter(Y_test, y_hat_sgd)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.grid()
plt.show()
```



```
In [77]: # Mean Squared Error - Fixed LR, Variable LR
MSE_flr_inbuilt = mean_squared_error(Y_test, y_hat_sgd)

print (f"Fixed LR MSE {MSE_flr_inbuilt}")
```

Fixed LR MSE 29.7729382260675

## Summarizing the results

```
In [78]: # A tabular representation of weights for all four modes
from prettytable import PrettyTable

x = PrettyTable()
x.field_names = ["Weights - Custom", "Weights - Inbuilt"]
```

```
In [79]: weights_sgd = sgd.coef_
```

```
In [80]: for i in range(12):
          x.add_row([fweights[0][i],weights_sgd[i]])
          print(x)
```

Weights - Custom	Weights - Inbuilt
-0.8951903669319022	-0.9957863796826973
1.129566519981019	1.3733331426243682
-0.35030211113827559	0.043633719299515004
1.139232115194231	0.3210297525187523
-1.3524086841272596	-1.6073743441292732
2.366946004315772	1.9590627259886977
-0.23322673440138084	-0.0808845687432791
-3.296705073959816	-3.231395894872585
2.012431578586961	2.8169369696694315
-1.3330193563306263	-2.268061926237042
-1.6670794991024582	-1.752852491026695
0.7388630842189096	0.9634424018569003

```
In [81]: # Custom SGD v/s SGD sklearn implementation
print('MSE of manual implementation = ',MSE_flr)
print('-'*50)
print('MSE of SGD sklearn implementation = ',MSE_flr_inbuilt)
```

```
MSE of manual implementation = 30.64473411784173
```

```
-----
MSE of SGD sklearn implementation = 29.7729382260675
```

## Summary

When compared manual SGD implementation with that of third party module of sklearn we pretty much see the same weights and MSE